


# class07\_machine\_learning

Xueran Zou

## Example of K-means clustering

The first step is to make up some data with a known structure, so we know what the answer should be.

```
tmp <- c(rnorm(30, mean = -3), rnorm(30, mean = 3))  
x <- cbind(x = tmp, y = rev(tmp))  
plot(x)
```



Now we have some structured data in `x`. Let's see if K-means is able to identify the two groups.

```
k <- kmeans(x, centers = 2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1  3.377489 -2.838476
2 -2.838476  3.377489
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 63.7635 63.7635
(between_SS / total_SS = 90.1 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Let's explore k.

```
k$size
```

```
[1] 30 30
```


```
k$centers
```

```
      x      y
1  3.377489 -2.838476
2 -2.838476  3.377489
```

```
k$cluster
```


```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

```
plot(x, col = k$cluster)
```




Now we can add the clusters centers:

```
plot(x, col = k$cluster)  
points(k$centers, col = 'blue', pch = 15)
```



An example when we select the wrong number of clusters for k-means

```
k_3 <- kmeans(x, centers = 3, nstart = 20)
plot(x, col = k_3$cluster)
```



## Example of Hierarchical Clustering

Let's use the same data as before, which we stored in "x". We will use the "hclust()" function.

```
clustering <- hclust(dist(x))  
clustering
```

Call:

```
hclust(d = dist(x))
```


Cluster method : complete

Distance : euclidean

Number of objects: 60

```
plot(clustering)
```

## Cluster Dendrogram




dist(x)  
hclust (\*, "complete")

Let's add a horizontal line

```
plot(clustering)
abline(h = 10, col = 'pink')
```

## Cluster Dendrogram



dist(x)  
hclust (\*, "complete")


To get our results (i.e., membership vector) we need to “cut” the tree. The function for doing that is `cutree()`

```
subgroups <- cutree(clustering, h = 20)
subgroups
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Plot this:

```
plot(x, col = subgroups)
```



You can also “cut” your tree with the number of clusters you want:

```
cutree(clustering, k = 2)
```

[illegible]

## Principal Component Analysis (PCA)

## Data import

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
head(x)
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93



5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

**Q1.** How many rows and columns are in your new data frame named `x`? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  5
```

```
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

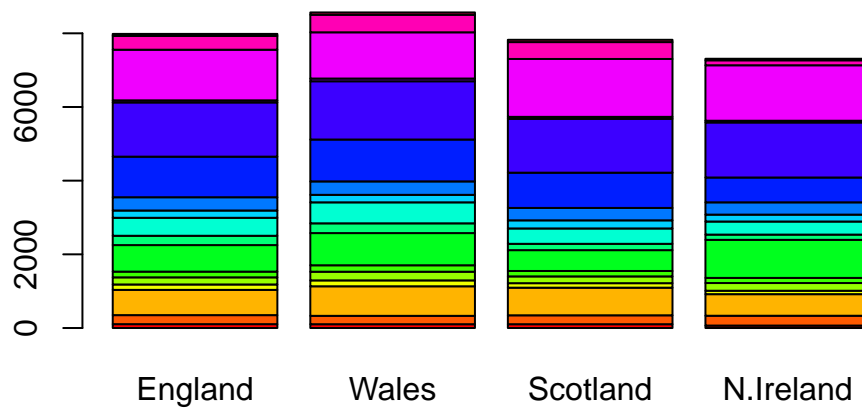
**Q2.** Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Using argument setting `row.names=1` can be more robust. Because when running the first approach (`x <- x[,-1]`) multiple times, it will constantly remove the first column of the data frame.

Now we can generate some basic visualizations

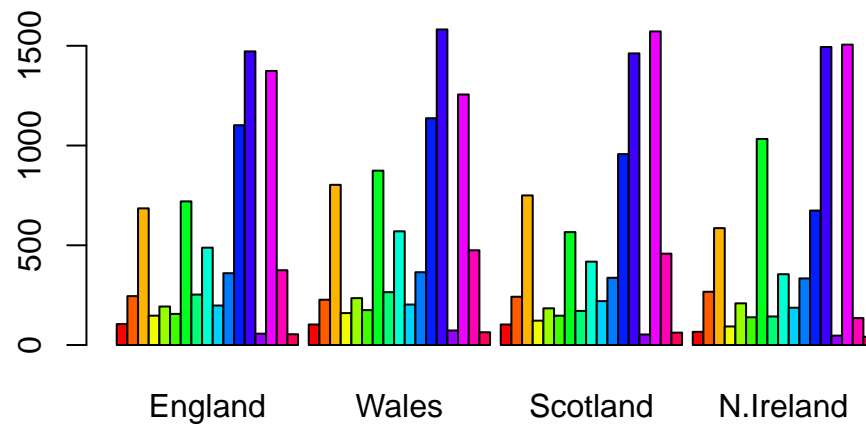
**Q3:** Changing what optional argument in the above **barplot()** function results in the following plot?

```
barplot(as.matrix(x), col = rainbow(nrow(x)))
```



Let's refine our barplot.


```
barplot(as.matrix(x), col = rainbow(nrow(x)), beside = TRUE)
```



**Q5:** Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Other visualization that can be useful:

```
pairs(x, col = rainbow(nrow(x)), pch = 15)
```



**Q6.** What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland takes less cheese, fish, fresh vegetables, fresh fruit and alcoholic drinks than the other countries, while it takes more fresh potatoes.

Let's apply PCA (Principal Component Analysis). For that, we need to use the command `prcomp()`. This function expects the transpose of our data.


```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's plot the PCA results:

```
plot(pca)
```



We need to access the results of the PCA.

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

```
[1] "prcomp"
```


We can explore the `pca$x` dataframe:

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13


Plot:

```
plot(x = pca$x[,1], y = pca$x[,2])
```




**Q7.** Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))  
text(pca$x[,1], pca$x[,2], colnames(x))
```



**Q8.** Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(x = pca$x[,1], y = pca$x[,2])
colors_countries <- c('orange', 'pink', 'blue', 'green')
text(x = pca$x[,1], y = pca$x[,2], colnames(x), col = colors_countries)
```



Calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 67 29 4 0
```

```
z <- summary(pca)
z$importance
```


	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	4.188568e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```






```
# Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



**Q9:** Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



## PCA of RNA-seq data

First step as always is loading the data:

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

**Q10:** How many genes and samples are in this data set?

```
dim(rna.data)
```

```
[1] 100 10
```

There are 100 genes and 10 samples in this data set.

Let's apply PCA:

```
pca <- prcomp(t(rna.data), scale=TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.348e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

```
plot(pca, main="Quick scree plot")
```




```
# Variance captured per PC
pca.var <- pca$sdev^2

# Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```


```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```



Let's plot the principal components 1 and 2.

```
plot(pca$x[,1], pca$x[,2], xlab = 'PC1', ylab = 'PC2')
```




```
colnames(rna.data)
```

```
[1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```


```
cols_samples <- c(rep('blue', 5), rep('red', 5))  
cols_samples
```

```
[1] "blue" "blue" "blue" "blue" "blue" "red" "red" "red" "red" "red"
```

```
plot(pca$x[,1], pca$x[,2], xlab = 'PC1', ylab = 'PC2', col = cols_samples)
```



```
barplot(pca$rotation[,1])
```



```
sort(pca$rotation[,1])
```

gene56	gene18	gene3	gene39	gene50	gene11
-0.103783479	-0.103774699	-0.103761385	-0.103744482	-0.103743341	-0.103719665
gene57	gene91	gene1	gene79	gene59	gene75
-0.103703675	-0.103698408	-0.103666005	-0.103639415	-0.103607438	-0.103592371
gene54	gene44	gene58	gene82	gene87	gene13
-0.103584153	-0.103504699	-0.103503980	-0.103481127	-0.103448562	-0.103399591
gene19	gene27	gene61	gene25	gene51	gene53
-0.103390599	-0.103374849	-0.103308945	-0.103302326	-0.103265591	-0.103245619
gene49	gene17	gene29	gene94	gene86	gene40
-0.103188532	-0.103013773	-0.102739689	-0.102692869	-0.102122719	-0.102003831
gene12	gene31	gene70	gene32	gene81	gene78
-0.102001924	-0.101768804	-0.101365212	-0.100677376	-0.100659777	-0.100499426
gene42	gene24	gene77	gene96	gene46	gene65
-0.098746675	-0.098284250	-0.097473626	-0.096658194	-0.096571619	-0.094219475
gene80	gene43	gene26	gene9	gene64	gene84
-0.093476477	-0.092001819	-0.085745836	-0.085460936	-0.078643996	-0.009263882
gene69	gene83	gene4	gene5	gene92	gene71
0.048197107	0.066065263	0.075320862	0.087428334	0.092534408	0.095664760
gene88	gene6	gene15	gene89	gene37	gene8
0.098226585	0.099670829	0.099993193	0.100038548	0.100467583	0.100759370
gene97	gene63	gene74	gene73	gene38	gene95
0.100787961	0.101468649	0.101747637	0.102001050	0.102080752	0.102142492
gene72	gene35	gene14	gene52	gene22	gene93
0.102347342	0.102382706	0.102478762	0.102519795	0.102725125	0.102950950
gene30	gene20	gene36	gene67	gene47	gene76
0.103044435	0.103121803	0.103412422	0.103453646	0.103502386	0.103514464
gene2	gene34	gene33	gene16	gene7	gene28
0.103514749	0.103525731	0.103592988	0.103598474	0.103609009	0.103638752
gene99	gene23	gene48	gene55	gene85	gene62
0.103649598	0.103681565	0.103682769	0.103695870	0.103698370	0.103713893
gene41	gene90	gene10	gene21	gene60	gene98
0.103716818	0.103777744	0.103783379	0.103787935	0.103805515	0.103837190
gene68	gene45	gene66	gene100		
0.103839510	0.103840183	0.103845454	0.103870820		




## Using ggplot

```
library(ggplot2)

df <- as.data.frame(pca$x)

ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```




Make it nicer.

```
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)


p
```



```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
theme_bw()
```

## PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data

## Gene loadings

Find the top 10 measurements (genes) that contribute most to PC1 in either direction.

```
loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes
```

```
[1] "gene100" "gene66" "gene45" "gene68" "gene98" "gene60" "gene21"
[8] "gene56" "gene10" "gene90"
```